Dynamic Qubit Allocation for Adversarial Quantum Routing:

From Single-Algorithm Evaluation to Multi-Allocator Framework

Comprehensive Analysis Across 60 Experimental Conditions

Piter Garcia
Graduate Assistant, AI/Quantum Computing
Rochester Institute of Technology
pzg8794@g.rit.edu

November 11, 2025
Version 2.0 - Enhanced Integration Report

Abstract

This report presents a comprehensive two-phase evaluation of quantum entanglement routing algorithms under adversarial conditions. **Phase 1** established baseline performance of EXPNeuralUCB across five attack scenarios (Stochastic, Markov, Adaptive, OnlineAdaptive, Baseline), demonstrating 74.4% Oracle efficiency with $<3\%$ degradation. **Phase 2** expanded evaluation to include Pursuit-based neural bandits (CPursuitNeuralUCB, iCPursuitNeuralUCB) and dynamic qubit allocation strategies (Fixed, Dynamic UCB, Random, Thompson Sampling), testing 60 total conditions (4 algorithms × 5 scenarios × 4 allocators × 3 horizons).

Key Findings: (1) **Pursuit algorithms dominate:** CPursuitNeuralUCB and iCPursuitNeuralUCB achieve 89% average efficiency (60%, 45% win rates) vs EXPNeuralUCB 78% (5% win rate). (2) **Thompson Sampling optimal:** 88.2% efficiency with 3× speedup. (3) **Fixed Allocator Trap discovered:** Single-allocator evaluation creates 6.4% generalization penalties. (4) **EXPNeuralUCB baseline validated:** Phase 1 results remain robust across allocators (74.4% avg), providing critical benchmark. (5) **Stochastic challenge confirmed:** Both phases identify pure randomness as hardest condition (73.2% EXP Phase 1, 77.4% Avg Phase 2).

Methodological Contribution: We introduce multi-allocator evaluation framework revealing that Phase 1's Fixed-only testing (standard practice) systematically favors specialists. Phase 2 demonstrates 17.3% performance improvement through algorithm evolution (EXP→Pursuit) and 3× speedup through allocator optimization (Fixed→Thompson).

Production Recommendations: Deploy iCPursuitNeuralUCB + Thompson (88-92% efficiency, CV $<8\%$) as primary. Use EXPNeuralUCB + Fixed as conservative baseline (74.4%, proven robustness from Phase 1).

Keywords: Quantum networking, Multi-armed bandits, Thompson Sampling, Adversarial robustness, Pursuit learning, Dynamic allocation

Contents

1 Phase 1: Baseline Evaluation Framework **3**
1.1 Original Research Context . . . 3
1.2 Phase 1 Methodology . . . 3
1.2.1 Attack Model Taxonomy . . . 3
1.3 Phase 1 Comprehensive Results . . . 3
1.3.1 Multi-Horizon Performance . . . 3
1.3.2 Algorithm Comparison . . . 3
1.4 Phase 1 Framework Validation . . . 4
1.4.1 Robustness Assessment . . . 4
1.5 Phase 1 Contributions . . . 4

2 Phase 2: Multi-Allocator Framework Expansion **4**
2.1 Expanded Methodology . . . 4
2.1.1 Algorithm Suite . . . 4
2.1.2 Allocator Comparison . . . 5
2.2 Phase 2 Comprehensive Results . . . 5
2.2.1 Per-Scenario Performance . . . 5
2.3 The Fixed Allocator Trap (Connecting Phases) . . . 5
2.3.1 Phase 1 Hidden Limitation . . . 5
2.3.2 Phase 2 Validation . . . 5
2.4 Attack-Allocator Interactions . . . 5
2.4.1 Synergies Discovered . . . 5
2.4.2 Conflicts Identified . . . 6

3 Integrated Discussion **6**
3.1 Why Pursuit Beats EXP (Cross-Phase Analysis) . . . 6
3.1.1 Phase 1 Success, Phase 2 Limitation . . . 6
3.1.2 Pursuit Superiority . . . 6
3.2 Capacity Configuration Impact . . . 6
3.3 Stochastic Challenge (Cross-Phase Validation) . . . 6

4 Production Deployment Guide **7**
4.1 Configuration Hierarchy . . . 7
4.2 Threat-Based Selection . . . 7

5 Conclusions **7**
5.1 Two-Phase Achievements . . . 7
5.2 Methodological Impact . . . 7
5.3 Performance Evolution . . . 8
5.4 Future Directions . . . 8

6 Acknowledgments **8**

1 Phase 1: Baseline Evaluation Framework

1.1 Original Research Context

Phase 1 established the foundational evaluation framework for adversarial quantum multi-armed bandit algorithms under systematic attack conditions, focusing on EXPNeuralUCB algorithm validation across five environmental conditions.

Phase 1 Key Results Summary

- **EXPNeuralUCB Superiority**: Consistent winner across all 5 environments with 17.4-18.3% gap over baselines
- **Stochastic Challenge**: Most difficult condition with 73.2% efficiency (2.5% degradation)
- **Robustness Score**: 9.7/10 rating with maximum 2.5% degradation under attacks
- **Oracle Efficiency**: 73.8-75.1% efficiency maintained (avg 74.4%)
- **Statistical Validation**: All differences significant at $p < 0.001$

1.2 Phase 1 Methodology

1.2.1 Attack Model Taxonomy

Phase 1 introduced graduated threat assessment across five environments:

Table 1: Phase 1 Attack Taxonomy and Difficulty Rankings

Environment	Characteristics	Oracle Gap	Difficulty
None	Deterministic optimal	**24.9%**	Easiest
Adaptive	Feedback-driven strategic	**24.9%**	2nd
OnlineAdaptive	Real-time adaptive	**25.4%**	3rd
Markov	Memory-dependent	**25.8%**	4th
Stochastic	Uniform random	**26.8%**	**Hardest**

Counterintuitive Discovery: Pure randomness (Stochastic) proved harder than sophisticated strategic attacks, establishing baseline challenge metric.

1.3 Phase 1 Comprehensive Results

1.3.1 Multi-Horizon Performance

Table 2: Phase 1 EXPNeuralUCB Learning Dynamics (Fixed Allocation: 8,10,8,9)

Scenario	4K Frames		6K Frames		8K Frames		Improve
	Reward	Gap	Reward	Gap	Reward	Gap	(pp)
None	3,702	41.2	4,213	33.1	**4,729**	**24.9**	**-16.3**
Stochastic	3,520	44.1	4,067	35.4	**4,612**	**26.8**	-17.3
Markov	3,621	42.5	4,139	34.3	4,673	25.8	-16.7
Adaptive	3,702	41.2	4,213	33.1	4,729	24.9	-16.3
OnlineAdaptive	3,676	41.7	4,183	33.6	4,701	25.4	-16.3
Average	**3,644**	**42.1**	**4,163**	**33.9**	**4,689**	**25.6**	**-16.5**

Phase 1 Insight: EXPNeuralUCB achieves consistent 74-75% Oracle efficiency with robust 16.5 pp improvement over 8K frames, validating theoretical regret bounds.

1.3.2 Algorithm Comparison

Statistical Validation: All gaps significant ($p < 0.001$), establishing EXPNeuralUCB as Phase 1 benchmark.

Table 3: Phase 1 Baseline Algorithm Comparison (8K frames, Fixed allocation)

Algorithm	Avg Reward	Oracle Eff	Gap vs EXP	Consistency
EXPNeuralUCB	**4,689**	**74.4%**	-	**0.8% CV**
NeuralUCB	3,580	56.9%	-17.4%	0.6% CV
EXPUCB	3,058	48.6%	-25.8%	0.5% CV
Oracle	6,298	100%	-	0% CV

1.4 Phase 1 Framework Validation

1.4.1 Robustness Assessment

Table 4: Phase 1 Robustness Scoring (Fixed Allocation)

Algorithm	Avg Degrad	Max Degrad	Recovery	Score	Grade
EXPNeuralUCB	**0.9%**	**2.5%**	**Fast**	**9.7/10**	**A+**
NeuralUCB	0.9%	2.6%	Medium	9.1/10	A
EXPUCB	1.0%	2.9%	Slow	8.9/10	A-

Key Validation: $<$3% degradation under worst attacks (Stochastic) establishes production-ready robustness threshold.

1.5 Phase 1 Contributions

Methodological:

- First systematic framework distinguishing stochastic failures from strategic attacks
- Graduated threat assessment (5 environments)
- Quantitative robustness benchmarking protocol

Empirical:

- EXPNeuralUCB validated: 74.4% avg efficiency, 9.7/10 robustness
- Stochastic identified as hardest challenge (counterintuitive)
- Performance benchmarks: 73.2% (worst) to 75.1% (best)

Phase 1 → Phase 2 Transition

Phase 1 Limitation Identified: Fixed allocation only (8,10,8,9) standard in literature but unrealistic for production deployment requiring dynamic resource management.
Phase 2 Motivation: Real quantum networks need adaptive allocation for variable demand, decoherence, topology changes. Phase 1's Fixed-only evaluation cannot predict performance under dynamic conditions.
Research Question: Does EXPNeuralUCB's robustness (Phase 1) generalize to diverse allocators, or is it Fixed-specific optimization?

2 Phase 2: Multi-Allocator Framework Expansion

2.1 Expanded Methodology

Phase 2 introduces **systematic multi-allocator evaluation**: 4 algorithms × 5 scenarios × 4 allocators = 60 conditions.

2.1.1 Algorithm Suite

Critical Discovery: EXPNeuralUCB (Phase 1 champion at 74.4% Fixed-only) drops to 77.8% avg across allocators—still competitive but no longer dominant. Pursuit algorithms achieve 89% with superior generalization.

Table 5: Phase 2 Algorithm Comparison (All conditions)

Algorithm	Avg Eff	CV%	Win Rate	Floor
iCPursuitNeuralUCB	**89.1%**	**7.9%**	**45%**	**72.9%**
CPursuitNeuralUCB	**89.0%**	8.5%	**60%**	64.6%
GNeuralUCB	84.2%	9.3%	30%	65.3%
EXPNeuralUCB	**77.8%**	**10.9%**	**5%**	**54.0%**

Table 6: Phase 2 Qubit Allocation Strategy Performance

Strategy	Config	Efficiency	Time	Note
Fixed	(8,10,8,9) static	87.7%	11.62h	Phase 1 baseline
Dynamic UCB	UCB $\beta = 2.0$	85.3%	11.24h	Adapts but underperforms
Random	$\varepsilon = 1.0$	77.3%	11.65h	Validates importance
Thompson	**(9,9,9,8) Bayesian**	**88.2%**	**3.52h**	**Best + 3× speedup**

2.1.2 Allocator Comparison

Efficiency Breakthrough: Thompson Sampling achieves Phase 1-level efficiency (88.2% vs 87.7% Fixed) with **66-70% faster execution**.

2.2 Phase 2 Comprehensive Results

2.2.1 Per-Scenario Performance

Table 7: Phase 2 Complete Performance Matrix (Best performer per scenario-allocator highlighted)

Scenario	Allocator	iCPursuit	CPursuit	GNeural	EXP
Stochastic	Fixed	91.6	**93.3**	88.1	86.0
	Dynamic	89.2	87.4	86.9	74.6
	Random	85.0	78.8	77.3	70.0
	Thompson	**94.0**	90.5	86.9	83.6
Adaptive	Fixed	**97.4**	92.8	87.7	90.2
	Dynamic	88.8	80.0	83.9	75.9
	Random	72.9	68.1	65.0	54.0
	Thompson	90.6	86.5	85.8	75.0
Phase 1 EXP	**Fixed only**	-	-	-	**74.4**

Integration Insight: Phase 1's EXPNeuralUCB (74.4% Fixed) aligns with Phase 2 Stochastic/Fixed (86.0%, different scaling). Pursuit algorithms improve 17-19% over Phase 1 baseline.

2.3 The Fixed Allocator Trap (Connecting Phases)

2.3.1 Phase 1 Hidden Limitation

Critical Discovery: Phase 1's Fixed-only evaluation (standard practice) cannot detect 6.4% generalization penalty. EXPNeuralUCB optimized for Fixed $\rightarrow$ fails to generalize. Pursuit algorithms designed for robustness $\rightarrow$ maintain performance.

2.3.2 Phase 2 Validation

Correlation: $r = 0.304$ between Fixed performance and penalty ($p < 0.05$). Algorithms with high Fixed-only scores (like Phase 1 EXP) exhibit systematic generalization failures.

Recommendation: Test algorithms across ≥ 3 allocators. Phase 1 methodology (Fixed-only) insufficient for production prediction.

2.4 Attack-Allocator Interactions

2.4.1 Synergies Discovered

- **Markov + Dynamic:** 88.5% vs 85.8% Fixed (+2.7%) — structured attacks align with UCB exploration

Table 8: Fixed Allocator Trap: Generalization Penalties

Algorithm	Fixed Perf	Other Avg	Penalty	Phase
EXPNeuralUCB	**82.6%**	**76.2%**	**6.4%**	Phase 2
CPursuitNeuralUCB	91.0%	88.2%	2.8%	Phase 2
iCPursuitNeuralUCB	92.9%	87.6%	5.3%	Phase 2
Phase 1 EXP	**74.4%**	**Unknown**	**?**	Phase 1

- **OnlineAdaptive + Thompson:** 90.6% vs 85.0% Dynamic (+5.6%) — real-time learning benefits
- **Stochastic + Thompson (Phase 1 link):** 90.5% avg — addresses Phase 1's hardest challenge

2.4.2 Conflicts Identified

- **Adaptive + Random:** 63.7% vs 90.2% Fixed (-26.5%) — catastrophic collapse
- **Adaptive + Dynamic:** 82.1% vs 90.2% Fixed (-8.1%) — exploration leakage
- **Stochastic + Random:** 77.4% vs 90.5% Thompson (-13.1%) — compounded uncertainty

Design Principle: Match allocator to attack structure. Phase 1's Fixed worked well for all attacks (no leakage). Phase 2 reveals Fixed prevents strategy inference—explains Phase 1 success.

3 Integrated Discussion

3.1 Why Pursuit Beats EXP (Cross-Phase Analysis)

3.1.1 Phase 1 Success, Phase 2 Limitation

EXPNeuralUCB Mechanism: Exponential weighting $w_t(a) \propto \exp(\eta \sum r_s)$ creates fast adaptation but brittle convergence. Under Fixed allocation (Phase 1), predictable resources enable stable learning $\rightarrow$ 74.4% efficiency.

Phase 2 Failure: Dynamic/Random allocation disrupts exponential convergence $\rightarrow$ 6.4% penalty. EXP's multiplicative updates amplify allocation noise.

3.1.2 Pursuit Superiority

Mechanism: Pursuit updates $p_{t+1}(a) = p_t(a) + \beta(\mathbf{1}_{a=a^*} - p_t(a))$ provide *adaptive damping*. $\beta < 1$ bounds updates ($\Delta p \leq \beta$), preventing overreaction to allocation variability.

Evidence:

- Pursuit CV 7.9-8.5% vs EXP 10.9% (more stable)
- Pursuit floor 64.6-72.9% vs EXP 54.0% (graceful degradation)
- Pursuit 2.8% penalty vs EXP 6.4% (better generalization)

3.2 Capacity Configuration Impact

Phase 1: Fixed 8,10,8,9 (35 total) with implied $T \approx 10K$ capacity **Phase 2:** $SC = 2T$ (8K capacity @ 4K frames) = *higher scarcity*

Interpretation: Phase 2's tighter constraints make 89% Pursuit efficiency more impressive than Phase 1's 74.4% EXP. Pursuit succeeds despite harder resource pressure.

3.3 Stochastic Challenge (Cross-Phase Validation)

Phase 1 Discovery: Stochastic hardest (73.2% EXP, 26.8% gap) **Phase 2 Validation:** Stochastic remains challenging:

- Stochastic/Random: 77.4% avg (worst combo)
- Stochastic/Thompson: 90.5% avg (best allocator for randomness)

Breakthrough: Thompson Sampling (Phase 2) specifically solves Phase 1's hardest problem—probabilistic allocation matches stochastic uncertainty.

4 Production Deployment Guide

4.1 Configuration Hierarchy

Tier 1 (Recommended): iCPursuitNeuralUCB + Thompson

- Performance: 88-92% efficiency, CV $< 8\%$
- Speed: 3.52h (3× faster than Phase 1 baseline)
- Robustness: 72.9% floor (graceful degradation)
- Use case: Production deployment, variable conditions

Tier 2 (Conservative): EXPNeuralUCB + Fixed

- Performance: 74.4% efficiency (Phase 1 validated)
- Speed: 11.62h (standard)
- Robustness: 9.7/10 score, $<3\%$ degradation
- Use case: Risk-averse, proven baseline

Tier 3 (Specialist): CPursuitNeuralUCB + Fixed

- Performance: 91.0% efficiency (highest)
- Win rate: 60% (most scenarios)
- Use case: Peak performance when Fixed feasible

4.2 Threat-Based Selection

Table 9: Allocator Selection by Attack Type

Attack	Allocator	Rationale
Stochastic	Thompson	Probabilistic matching (90.5% avg, solves Phase 1 challenge)
Markov	Dynamic	Structure exploitation (+2.7% synergy)
Adaptive	Fixed	Prevents strategy leakage (6.4% penalty avoided)
OnlineAdaptive	Thompson	Real-time adaptation (+5.6% synergy)
Unknown	Thompson	Best generalization (88.2% avg)

5 Conclusions

5.1 Two-Phase Achievements

Phase 1: Established EXPNeuralUCB baseline (74.4% efficiency, 9.7/10 robustness) under Fixed allocation. Validated theoretical robustness, identified Stochastic as hardest challenge, created reproducible evaluation framework.

Phase 2: Expanded to multi-allocator, multi-algorithm evaluation (60 conditions). Discovered Fixed Allocator Trap (6.4% penalty), validated Pursuit superiority (89% vs 78%), achieved 3× speedup with Thompson, established production recommendations.

5.2 Methodological Impact

Fixed Allocator Trap: Phase 1's Fixed-only evaluation (standard in quantum networking literature) systematically favors specialists over generalists. Phase 2 reveals this creates deployment blindspots—algorithms succeed in testing but fail under production variability.

Framework Evolution: Phase 1 (5 scenarios, 1 allocator) $\rightarrow$ Phase 2 (5 scenarios, 4 allocators) represents 4× expansion in evaluation comprehensiveness. Multi-allocator testing should become standard practice.

5.3 Performance Evolution

Table 10: Cross-Phase Performance Summary

Metric	Phase 1 EXP	Phase 2 EXP	Phase 2 Pursuit	Improvement	Method
Avg Efficiency	74.4%	77.8%	**89.0%**	**+14.6 pp**	Algorithm
Execution Time	11.62h	11.62h	**3.52h**	**-8.1h**	Allocator
CV (Stability)	0.8%	10.9%	**7.9%**	Mixed	Both
Win Rate	100%	5%	**60%**	Variable	Context

Combined Gains: Algorithm evolution (EXP→Pursuit) + Allocator optimization (Fixed→Thompson) = 14.6 pp efficiency improvement + 3× speedup.

5.4 Future Directions

Phase 3 (iCMAB Integration): Incorporate predictive intelligence into Pursuit framework, building on Phase 1's robustness foundation and Phase 2's generalization insights.

Phase 4 (Hybrid Evaluation): Test EXP+iCMAB vs Pursuit across all allocators to determine if prediction compensates for EXP's generalization penalty.

Phase 5 (Theoretical Analysis): Derive regret bounds for Thompson+Pursuit. Conjecture: $O(\sqrt{T \log T})$ vs standard $O(T^{2/3})$ due to faster convergence demonstrated in Phase 2.

Interactive Framework Access

Phase 1 Framework (Original):
Dynamic Environment Test Framework
Phase 2 Framework (Enhanced):
Multi-Allocator Evaluation Suite (Available upon request)

6 Acknowledgments

This research conducted as Graduate Assistant work at RIT's AI/Quantum Computing track. Phase 1 completed September 2025, Phase 2 completed November 2025. Special thanks to Dr. Devroop Kar for iCMAB consultation.